

APNNS/ IEEE

Deep Learning and Artificial Intelligence

Summer School (DLAI3)

29 June – 3 July 2020



Doing Deep Learning: The hype and the reality

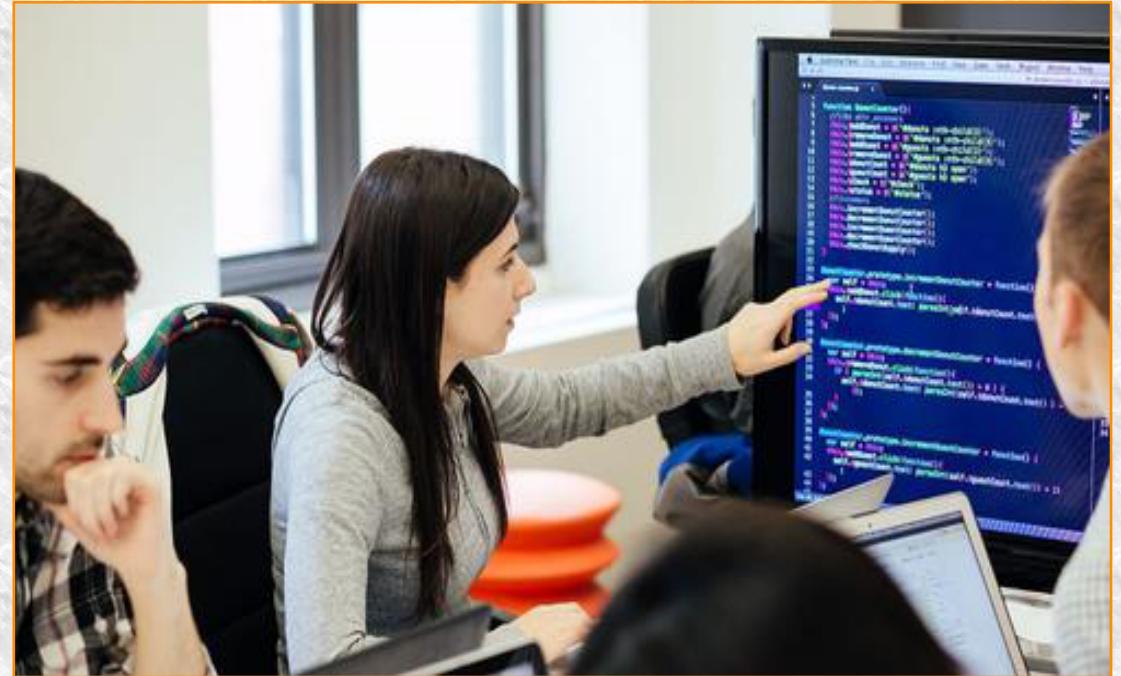
Sally E. Goldin
Department of Computer Engineering
King Mongkut's University of Technology Thonburi

The brave new world of deep learning



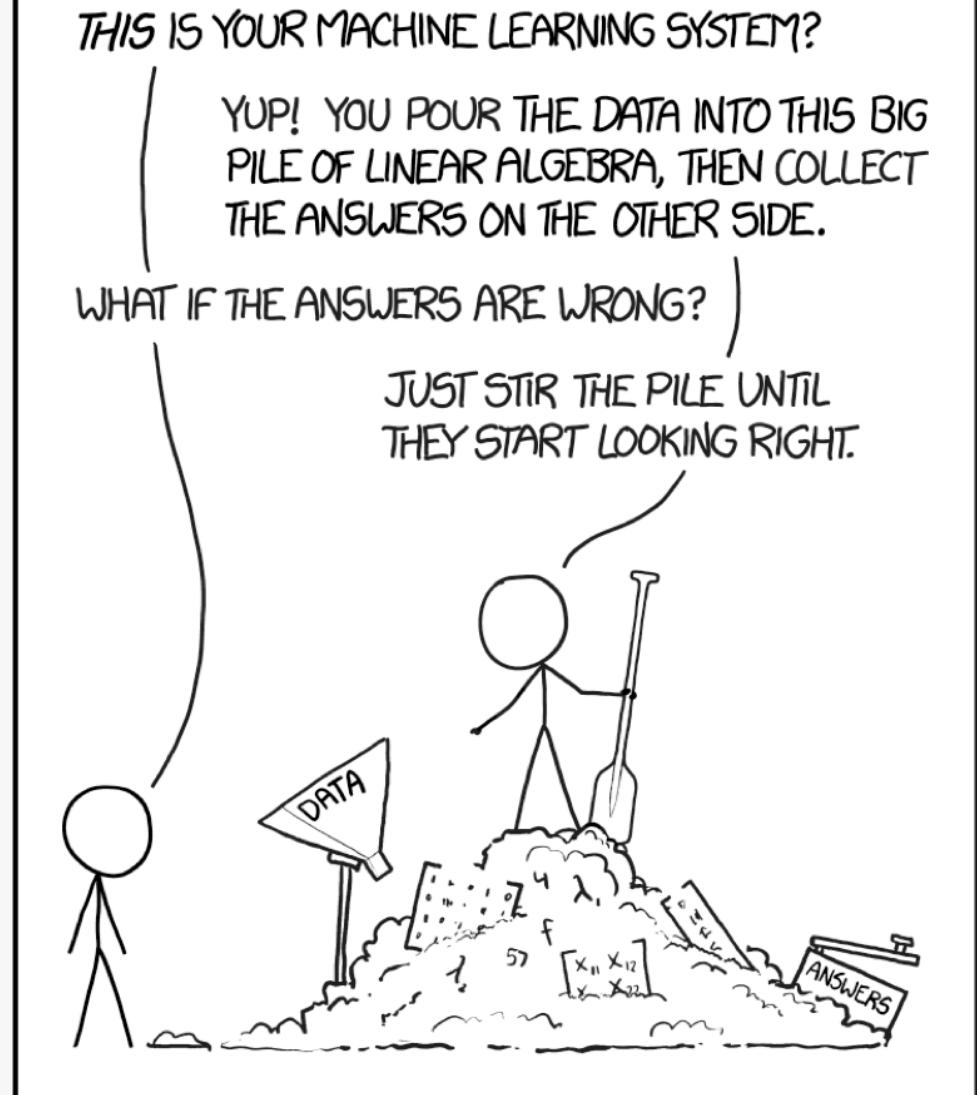
Accessible technologies and tools

- ***Hardware / infrastructure***
 - Cloud computing
 - GPGPU
- ***Software frameworks***
 - Tensorflow
 - PyTorch
 - Keras
 - Caffe
- ***Data sources***
 - Mobile devices
 - IOT
- ***Pre-trained models***



Anyone can build a DL model

Right?



About my background

PhD from Carnegie-Mellon University

During the first “golden age” of AI: Herb Simon, Allen Newell, Marvin Minsky, Patrick Winston, John McCarthy...



Research Scientist at the Rand Corporation

Software engineer, architect, project leader



Lecturer & researcher at KMUTT

*Two dozen programming languages
Many publications
1500+ students...*

Case study: My first real DL project

- *Successes*
- *Challenges*
- *Learning*

In collaboration with Professor Yuanjie Xiao
High Speed Railway Engineering Laboratory
Central South University, Changsha, China



The problem

Given a photo of a railway embankment under construction, estimate the level of compaction

Approach: Train a DL model to discriminate between different known gradations which can be mapped to estimates of compaction



Gradation 1

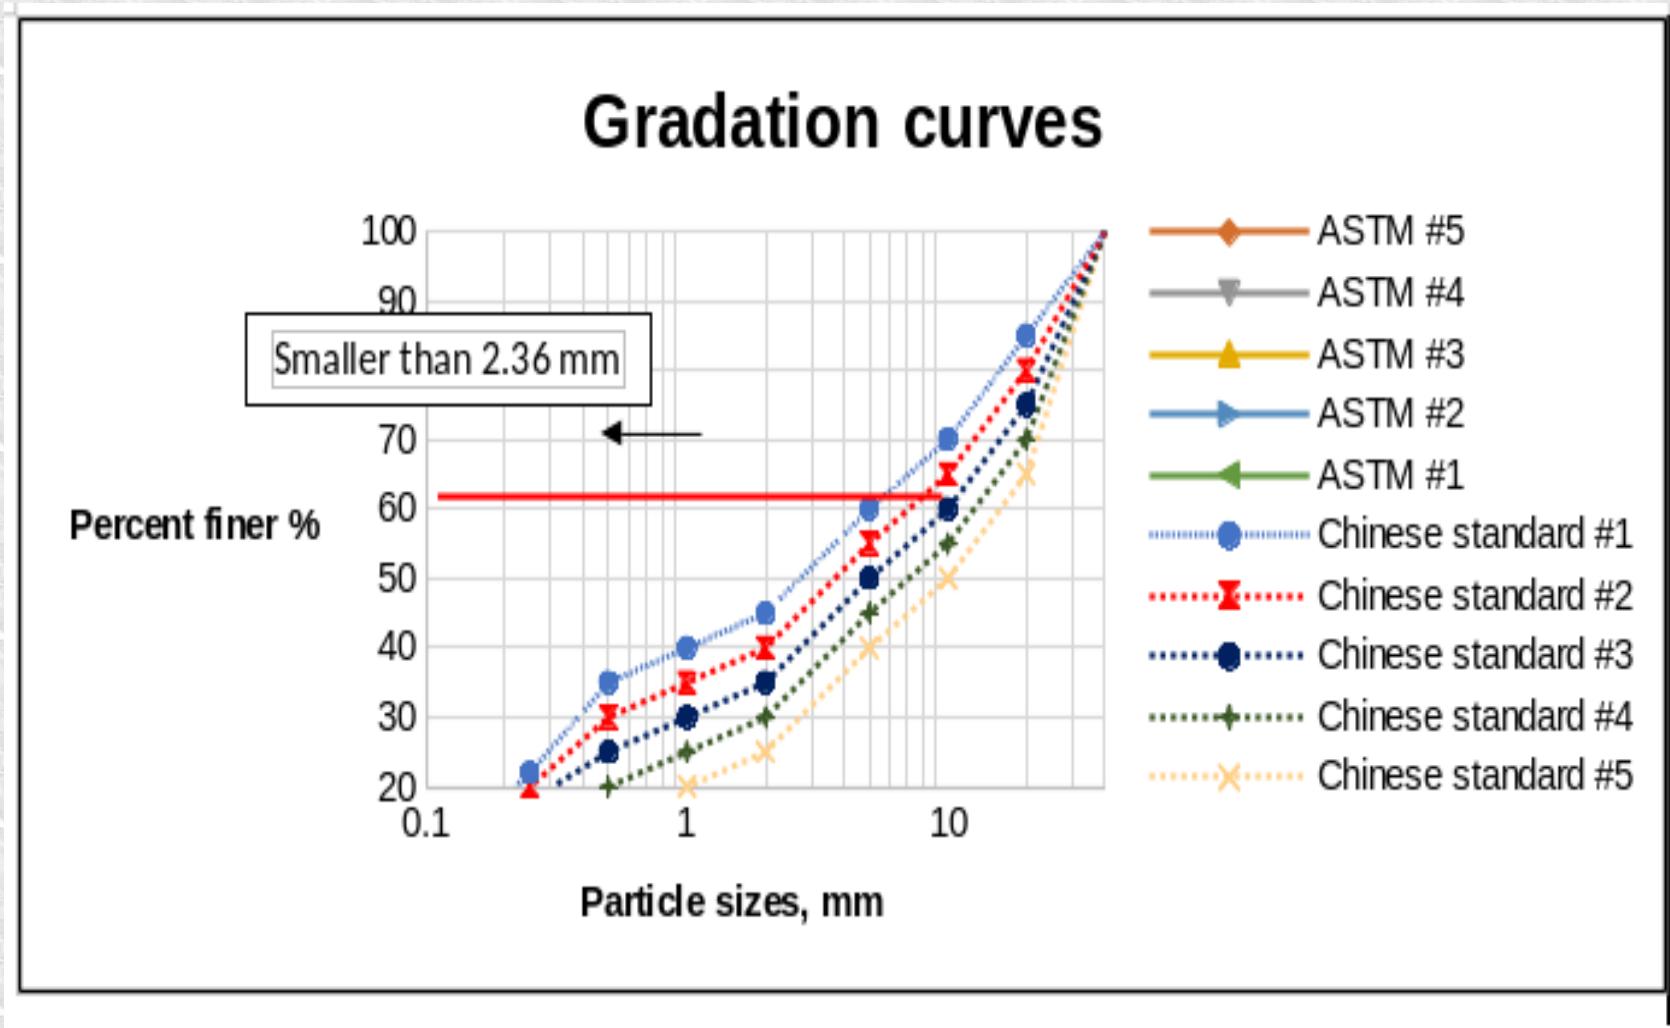


Gradation 3



Gradation 5

Different mixes of particle size

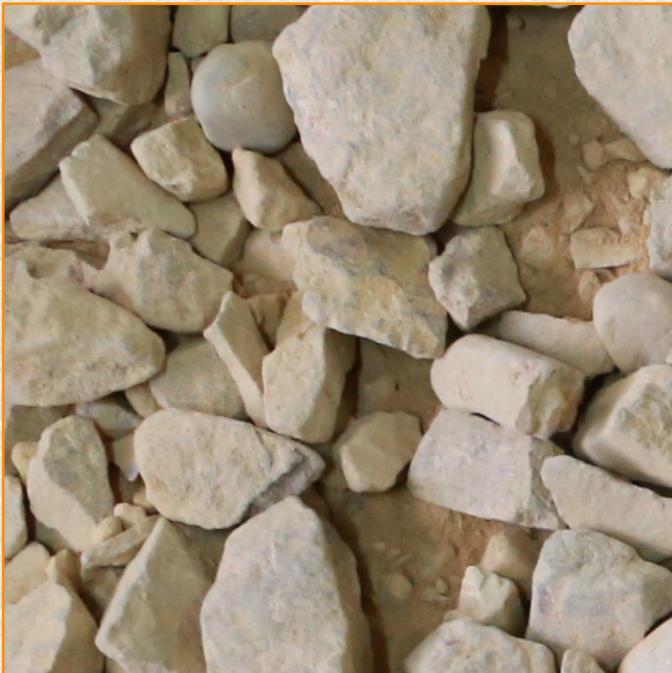


Challenge 1: Development environment

- Google Cloud: costs impossible to predict
- Configured dedicated virtual machine
 - Kubuntu 18
 - Python 3.7
 - PyTorch
 - No GPGPU
- PyTorch “getting started” tutorial crashed!
- Had to move VM from AMD to Intel-based host
- Continued to see other bugs (e.g. processes spontaneously killed)

Exploratory data set

- Only two categories, not standard gradations
- 200 images of each, 4896 x 3264 pixels – extracted multiple 500 x 500 subsets
- ~6300 training images per class, 1600 validation images



Class 1



Class 2

Best DL model

- **Model architecture**
 - 3 convolutional layers (6,16,32) with ReLU and max pooling
 - 3 fully connected layers (128,64,32) – 2 outputs
 - Cross-entropy loss error function, stochastic gradient descent back propagation, learning rate 0.001, momentum 0.9
- **Data loading / transformation**
 - Image load direct from disk
 - Gray scale conversion, sub-sample to 32 x 32 for input
 - Random 25% brightness and contrast “jitter”

Note: very similar to the generic model used in the PyTorch tutorials.

Success 1: Acceptable accuracy

- **92% accuracy** on validation images after 150 epochs
- Minimum loss still declining
- Probably could push accuracy higher
- Quite slow to train
 - About 5 minutes/epoch*
 - Total time 13 hours*
- These were not our real target classes

Standardized data set

- *Five categories*
- *Standard gradations*
- *Pre-processed like exploratory data set*
 - 200 images of each, 4896 x 3264
 - Extracted multiple 500 x 500 subsets
 - ~5600 training images per class, 1400 validation images
- *Total data volume approximately 10 GB (8GB training)*



Challenge 2: Learning speed

- Modified previous model (5 classes, 0.01 learning rate)
- 120 epochs -> loss 1.45, **accuracy 27%**
- Each epoch takes > 10 minutes, > 20 hours
- ***Experiments***
 - Bigger subsets (1000 x 1000) – more representative texture
 - Two gradations only (1 and 5) – easier discrimination – 240 epochs, loss 0.062, **accuracy 62%**
 - This took days! And still just two classes



New development environment

- Physical machine: AMD Ryzen 1850 MHz processor, 8 cores, 16 GB memory
- CentOS 8, Python 3.7.4
- PyTorch with CUDA 10.2 support
- Gigabyte RTX 2060 GPGPU
- 2 TB disk space

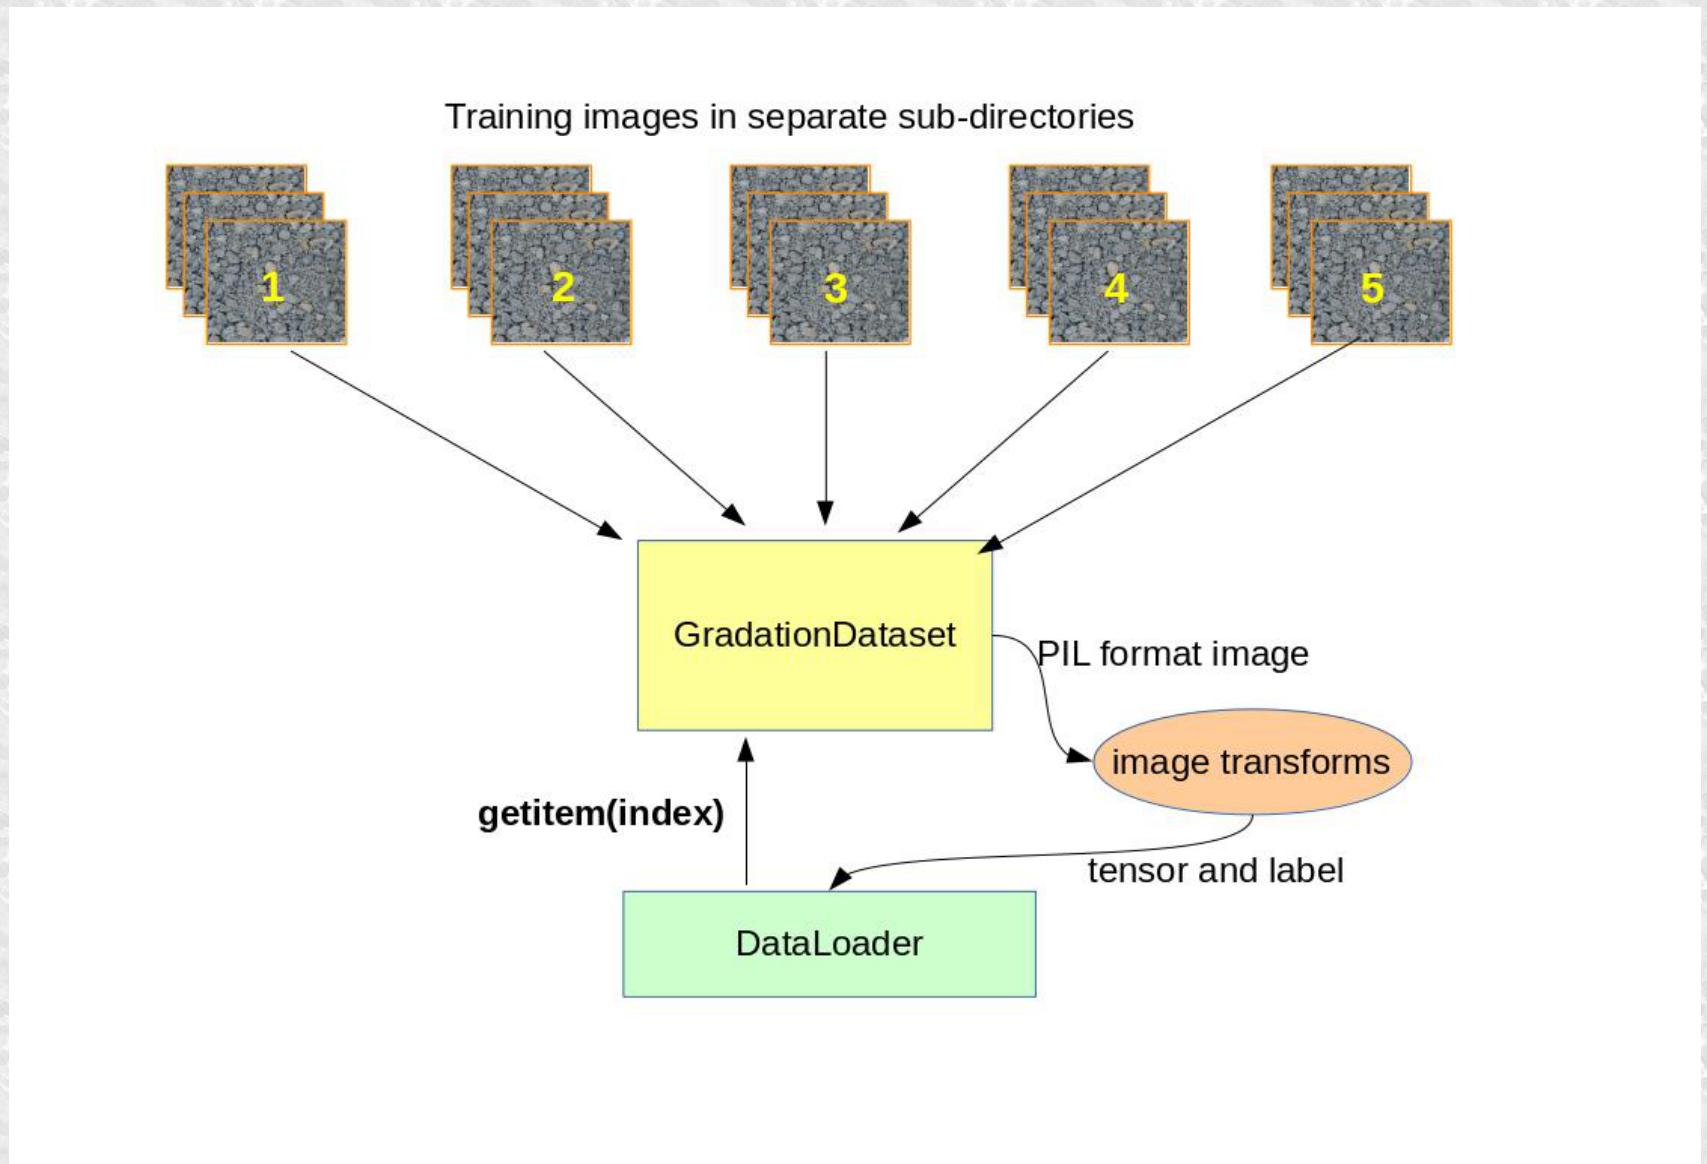
Added code to move model and data to the GPGPU device (cuda) during training



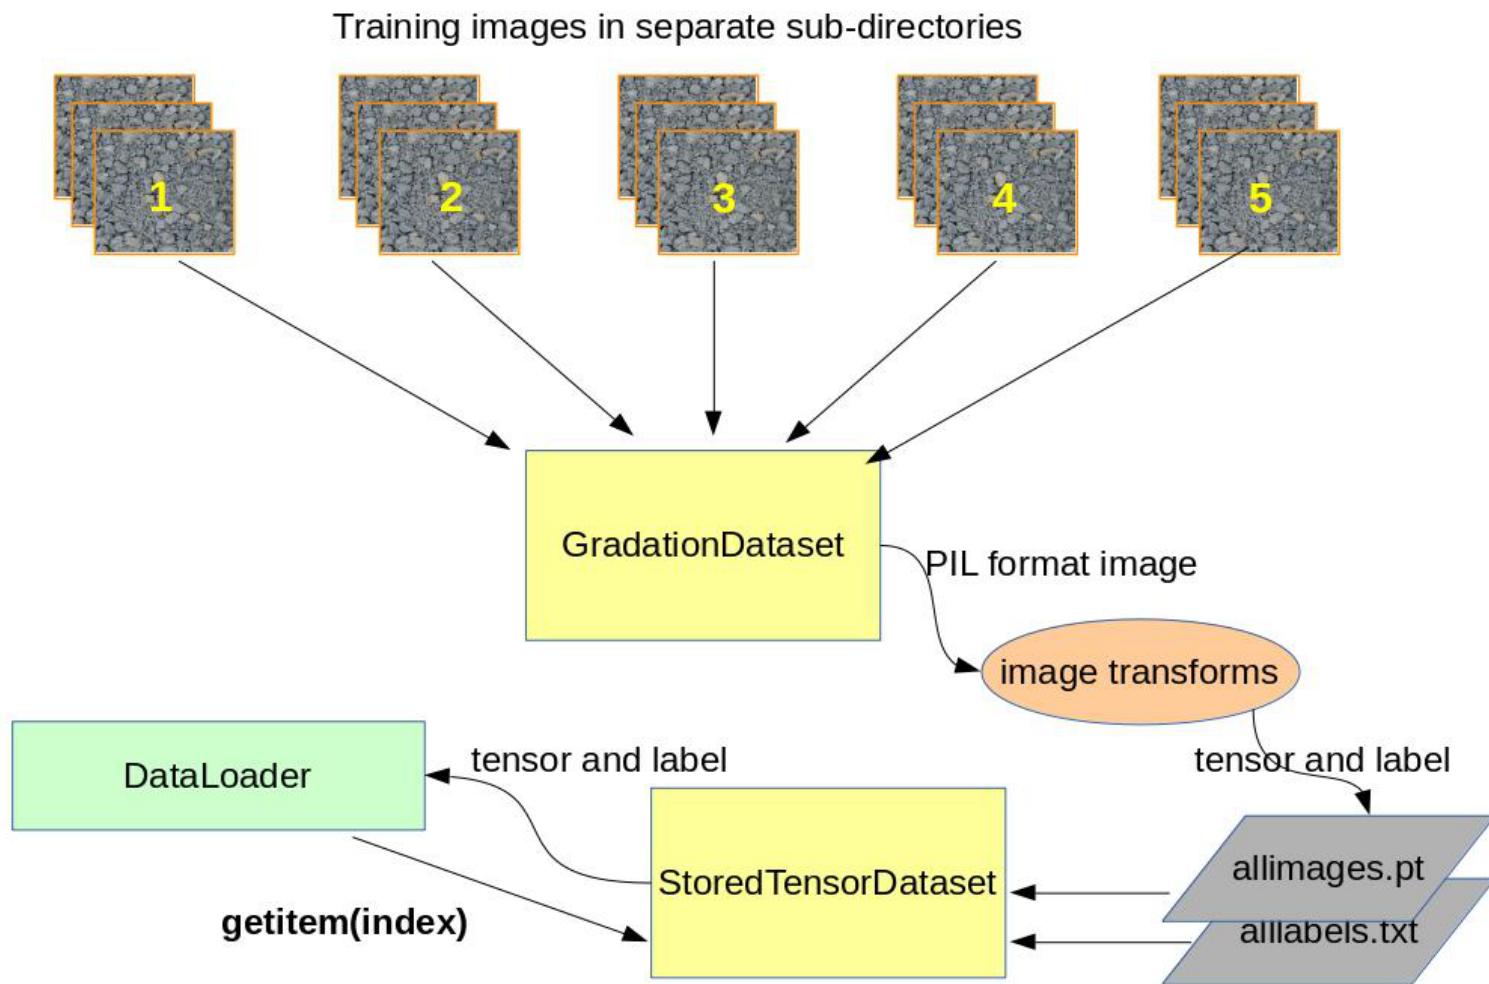
Challenge 3: GPGPU does not help

- Each epoch **still requires more than 9 minutes**
- 100 epochs, 15 hours total, loss down to only 1.203
- **Why?**
 - Sending all training data (8 GB) to the cuda device, each epoch
 - Data transfer time swamping computation time
- **Solution**
 - Pre-process (gray transform, subsample, color jitter) all training images and save as tensors in PyTorch serialization format
 - Create a new data loader class that loads all training data as tensors at the start of training
 - Each sample still sent to cuda device separately, but this is a small amount of data

Original data loading strategy



New data loading strategy

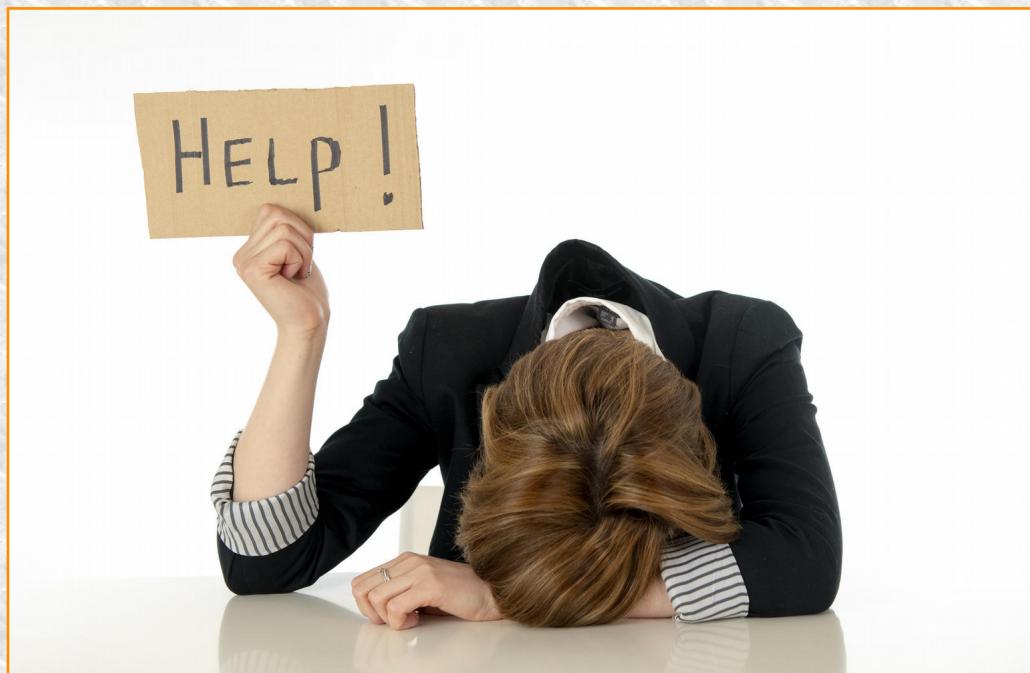


Success 2: Greatly improved speed

Each epoch now takes about 10 seconds

Challenge 4: Accuracy is still poor

1500 epochs – loss 0.657 – accuracy 30%



Analysis

- Is this a model problem?
- Or a data problem?
- Given the success this model had with the exploratory data, look at the data first
- ***New image subsets***
 - 2400 x 2400 (instead of 1000 x 1000) to better capture texture
 - Increased model input size from 32x32 to 64x64
 - Note training data images now require > 50 GB
 - However, we only need to pre-process this once, to create the serialized tensor data, which is about 120 MB



Success 3: Improved accuracy

- After 1100 epochs, loss 0.000, **accuracy 61%**
- Model is clearly over-fit (classifies the training data perfectly)
- Consulted PyTorch forum for suggestions
 - “Make sure you use Softmax as final layer”
 - “Try using dropouts”
 - “Add more training data”
 - “Optimize hyper-parameters”



More challenges...

- Added Softmax layer and tried dropouts
- Performance got much worse (no learning after 6000 epochs)
- Did not improve after removing the dropouts
- After some research I discovered that the Cross Entropy Loss class in PyTorch incorporates Softmax already

When you're a new to a field,
it's easy to be led astray!



Latest success

- *Doubled the number of training images*
 - 2600 total per class
 - More than 100 GB of raw data
- *Removed Softmax layer*
- *200 iterations, 0.000 loss, 90% accuracy*



More work to be done

- Increase accuracy (hyper-parameters...)
- Test against field images
- Integrate into real time environment



Lessons learned

- **PyTorch** – powerful, flexible, multiple levels of abstraction. well-supported – but not bug free
- **Data** may be more important than model structure
- Deep learning makes heavy demands on **hardware**, both computational power and disk space
- Deep learning requires time to configure the **software** environment
- **Detailed notes** on experiments are essential
- **Incorrect advice** can lead you astray

Deep learning is powerful...



but it's not magic!

Questions?



Contact information

Email: seg@goldin-rudahl.com

LinkedIn: <https://www.linkedin.com/in/sally-e-goldin-b77595b/>

Website: <http://windu.cpe.kmutt.ac.th>

YouTube: <https://tinyurl.com/youtubeseg>